

## EA's in practice

1

## Working with Evolutionary Algorithms

- Experiment design
- Algorithm design
- Test problems
- Measurements and statistics
- Some tips and summary

Slides adapted from chapter 9 of Eiben and Smith's book Introduction to EC  
<http://www.evolutionarycomputation.org/slides/>

2

## Experimentation

- Has a **goal** or goals
- Involves **algorithm** design and implementation
- Needs **problem(s)** to run the algorithm(s) on
- Amounts to **running** the algorithm(s) on the problem(s)
- Delivers **measurement data**, the results
- Is concluded with **evaluating** the results in the light of the given goal(s)
- Is often **documented**

3

## Experimentation: Goals

- Get a good solution for a given problem
- Show that EC is applicable in a (new) problem domain
- Show that *my EA* is better than *benchmark EA*
- Show that EAs outperform traditional algorithms (sic!)
- Find best setup for parameters of a given algorithm
- Understand algorithm behavior (e.g. pop dynamics)
- See how an EA scales-up with problem size
- See how performance is influenced by parameters
- ...

4

## Example: Production Perspective

- Optimising Internet shopping delivery route
  - Different destinations each day
  - Limited time to run algorithm each day
  - Must *always* be *reasonably* good route in limited time




5

## Example: Design Perspective

- Optimising spending on improvements to national road network
  - Total cost: billions of Euro
  - Computing costs negligible
  - Six months to run algorithm on hundreds computers
  - Many runs possible
  - Must produce *very* good result just *once*



6




## Perspectives of goals

- **Design** perspective:  
find a **very good** solution at least **once**
- **Production** perspective:  
find a **good** solution at **almost every run**
- **Publication** perspective:  
must meet **scientific standards** (huh?)
- **Application** perspective:  
good enough is **good enough** (verification!)

These perspectives have very different implications on evaluating the results (yet often left implicit)


7



## Algorithm design

- **Design a representation**
- **Design a way of evaluating an individual**
- **Design suitable mutation operator(s)**
- **Design suitable recombination operator(s)**
- **Decide how to select individuals to be parents**
- **Decide how to select individuals for the next generation (how to manage the population)**
- **Decide how to start: initialization method**
- **Decide how to stop: termination criterion**

8




## Test problems

- 5 DeJong functions
- 25 “hard” objective functions
- Frequently encountered or otherwise important variants of given practical problem
- Selection from recognized benchmark problem repository (e.g. TSP benchmark instances from the OR repository)
- Problem instances made by random generator

Choice has severe implications on

- generalizability and
- scope of the results


9



## Bad example (1/2)

- I invented “tricky mutation”
- Showed that it is a good idea by:
  - Running standard (?) GA and tricky GA
  - On 10 objective functions from the literature
  - Finding tricky GA better on 7, equal on 1, worse on 2 cases
- I wrote it down in a paper
- And it got published!
- Q: what did I learned from this experience?
- Q: is this good work?


10



## Bad example (2/2)

- What did I (my readers) not learn:
  - How **relevant** are these results (test functions)?
  - What is the **scope of claims** about the superiority of the tricky GA?
  - Is there a **property distinguishing** the 7 good and the 2 bad functions?
  - Are my results **generalizable**? (Is the tricky GA applicable for other problems? Which ones?)

11



## Getting Problem Instances (1/3)

- Testing on **real data**
- **Advantages:**
  - Results could be considered as very relevant viewed from the application domain (data supplier)
- **Disadvantages**
  - Can be over-complicated
  - Can be few available sets of real data
  - May be commercial sensitive – difficult to publish and to allow others to compare
  - Results are hard to generalize

12

## Getting Problem Instances (2/3)

- Standard data sets in problem repositories, e.g.:
  - OR-Library  
<http://www.ms.ic.ac.uk/info.html>
  - UCI Machine Learning Repository  
[www.ics.uci.edu/~mlearn/MLRepository.html](http://www.ics.uci.edu/~mlearn/MLRepository.html)
- Advantage:
  - Well-chosen problems and instances (hopefully)
  - Much other work on these → results comparable
- Disadvantage:
  - Might still miss crucial aspects
  - Algorithms get tuned for popular test suites

13

## Getting Problem Instances (3/3)

- Problem instance generators produce simulated data for given parameters, e.g.:
  - GA/EA Repository of Test Problem Generators  
<http://www.cs.uwo.edu/~wspears/generators.html>
- Advantage:
  - Allow very systematic comparisons for them
    - can produce many instances with the same characteristics
    - enable gradual traversal of a range of characteristics (hardness)
  - Can be shared allowing comparisons with other researchers
- Disadvantage
  - Not real – might miss crucial aspect
  - Given generator might have hidden bias

14

## Basic rules of experimentation

- EAs are stochastic →  
never draw any conclusion from a single run
  - perform sufficient number of independent runs
  - use statistical measures (averages, standard deviations)
  - use statistical tests to assess reliability of conclusions
- EA experimentation is about comparison →  
always do a fair competition
  - use the same amount of resources for the competitors
  - try different comp. limits
  - use the same performance measures

15

## Things to Measure

- Many different measures. Examples:
- Average result in given time
  - Average time for given result
  - Proportion of runs within % of target
  - Best result over  $n$  runs
  - Amount of computing required to reach target in given time with % confidence
  - ...

16

## What time units do we use?

- Elapsed time?
  - Depends on computer, network, etc...
- CPU Time?
  - Depends on skill of programmer, implementation, etc...
- Generations?
  - Difficult to compare when parameters like population size change
- Evaluations?
  - Evaluation time could depend on algorithm, e.g. direct vs. indirect representation

17

## Measures

- Performance measures (off-line)
  - Efficiency (alg. speed)
    - CPU time
    - No. of steps, i.e., generated points in the search space
  - Effectivity (alg. quality)
    - Success rate
    - Solution quality at termination
- “Working” measures (on-line)
  - Population distribution (genotypic)
  - Fitness distribution (phenotypic)
  - Improvements per time unit or per genetic operator
  - ...

18

## Performance measures

- No. of generated points in the search space  
= no. of fitness evaluations  
(don't use no. of generations!)
- **AES: average no. of evaluations to solution**
- **SR: success rate** = % of runs finding a solution (individual with acceptable quality / fitness)
- **MBF: mean best fitness** at termination, i.e., best per run, mean over a set of runs
- **SR  $\neq$  MBF**
  - Low SR, high MBF: good approximizer (more time could possibly help)
  - High SR, low MBF: "Murphy" algorithm (few very bad runs)

19

## Fair experiments

- **Basic rule: use the same computational limit for each competitor**
- Allow each EA the same no. of evaluations, but
  - Beware of hidden labour, e.g. in heuristic mutation operators
  - Beware of possibly fewer evaluations by smart operators
- EA vs. heuristic: allow the same no. of steps:
  - Defining "step" is crucial, might imply bias!
  - Scale-up comparisons alleviate this bias

20

## Better example: problem setting

- I invented myEA for problem X
- Looked and found 3 other EAs and a traditional benchmark heuristic for problem X in the literature
- Asked myself when and why is myEA better

21

## Better example: experiments

- Found/made problem instance generator for problem X with 2 parameters:
  - $n$  (problem size)
  - $k$  (some problem specific indicator)
- Selected 5 values for  $k$  and 5 values for  $n$
- Generated 100 problem instances for all combinations
- Executed all alg's on each instance 100 times (benchmark was also stochastic)
- Recorded AES, SR, MBF values w/ same comp. limit
- Put my program code and the instances on the Web

22

## Better example: evaluation

- Arranged results "in 3D" ( $n, k$ ) + performance  
(with special attention to the effect of  $n$ , as for scale-up)
- Assessed statistical significance of results
- Found the niche for my\_EA:
  - Weak in ... cases, strong in ... cases, comparable otherwise
  - Thereby I answered the "when question"
- Analyzed the specific features and the niches of each algorithm thus answering the "why question"
- Learned a lot about problem X and its solvers
- Achieved generalizable results, or at least claims with well-identified scope based on solid data
- Facilitated reproducing my results  $\rightarrow$  further research

23

## Some general tips

- **Be organized**
- Decide what you want & define appropriate measures
- Choose test problems carefully
- Make an experiment plan (estimate time when possible)
- Perform sufficient number of runs
- Keep all experimental data (never throw away anything)
- Use good statistics ("standard" tools from Web, MS, R)
- Present results well (figures, graphs, tables, ...)
- Watch the scope of your claims
- Aim at generalizable results
- Publish code for reproducibility of results (if applicable)
- Publish data for external validation (open science)

24